

Gebze Institute of Technology
Department of Computer Engineering
BIL 565 / 463
Computer Vision
Spring 2024

Semester Project

Title: Learning Keypoint Detectors and Descriptors using OpenCV Pseudo-Labels

Due Date: June 9th 2025

Objective:

Students will design, train, and evaluate a deep learning model to detect and describe keypoints in images. The keypoint labels will be generated using classical OpenCV methods. The trained system will be evaluated on publicly available RGB-D datasets and stereo correspondence tasks.

Task Description:

1. Dataset Collection

- Use RGB images from the [RGB-D Object Dataset](#) and the Middlebury stereo dataset.
- Focus on high-resolution, well-lit scenes with diverse object structures.

2. Pseudo-Label Generation

- Use OpenCV feature detectors (e.g., SIFT, ORB, FAST) to generate keypoint locations as training labels.
- Generate keypoint descriptors with OpenCV as ground truth for descriptor learning.

3. Network Design

- Design a CNN or ViT-based architecture that takes an image as input and produces:
 - A 2D keypoint probability heatmap
 - A 1D descriptor vector (e.g., 32-128 dimensions) per keypoint
- There should be a method to select keypoints from the heatmap with some threshold values

4. Training

- Train your model on RGB-D Object Dataset images using pseudo-labels.
- Augment data with synthetic views or transformations.
- Clearly state the loss functions and optimizer used.

5. Evaluation

- Evaluate the model on the Middlebury stereo dataset by using:
 - Keypoint matching on rectified stereo pairs
 - Descriptor matching (e.g., using cosine or L2 distance)
- Compare the correspondence results with classical OpenCV methods (SIFT, ORB, etc.)
- Provide numerical evaluations such as repeatability, matching accuracy, precision-recall.

6. Reporting

- Submit a **detailed report** that includes:
 - Network architecture and training strategy
 - Visualizations of predicted keypoints and descriptors
 - Quantitative performance metrics
 - Discussion of results, including failure cases
 - Projects **will not be graded** without a report.
-

Deliverables:

1. Source code (Python, PyTorch or TensorFlow recommended)
 2. Trained model weights
 3. A report (PDF, max 6 pages excluding references)
 4. Short demo video or interactive notebook
-

Performance Metrics Suggestions:

Metric	Description
Repeatability	Ratio of correctly repeated keypoints in stereo pairs
Matching Accuracy	% of correct keypoint matches
Descriptor Quality	Descriptor retrieval precision and recall (at fixed threshold)
Inference Speed	Time per image at test time

Additional Notes:

- You may use pre-trained CNN backbones if clearly documented.
- AI tools like ChatGPT or Copilot are allowed; include an appendix section describing how you used them.
- Collaboration rules: Discuss ideas, but **code and report must be written individually**.